Request length alone does not determine observed completion latency

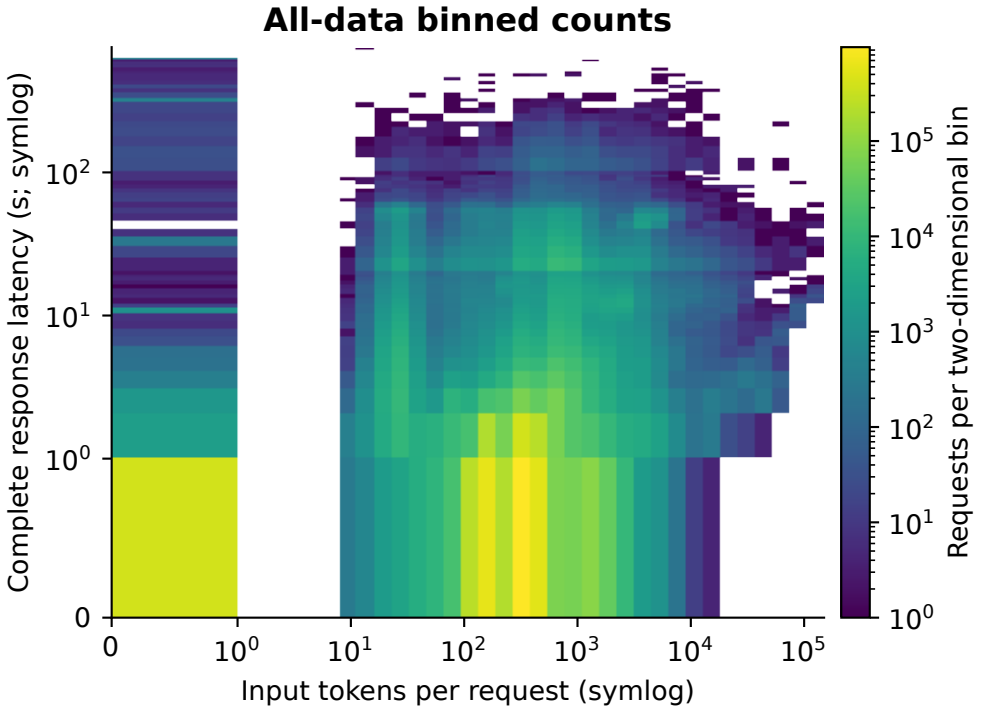

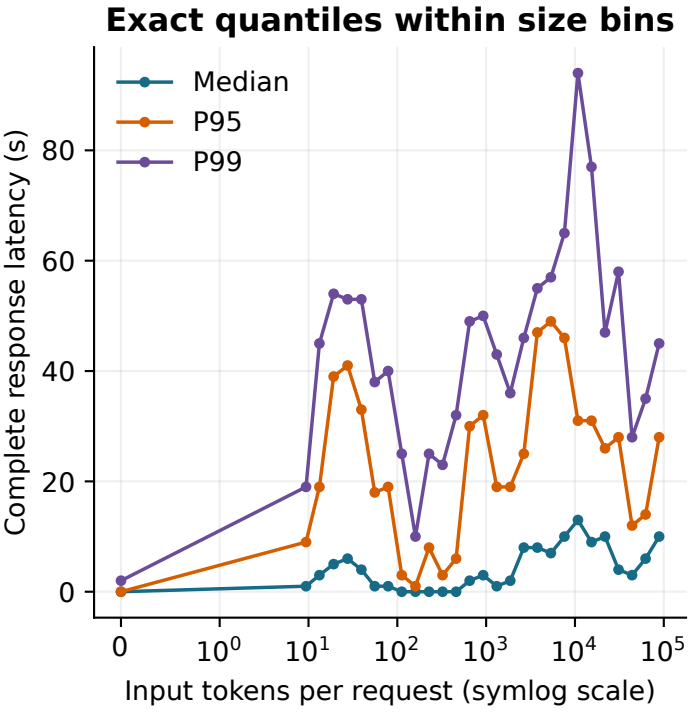

Source: HPMLL/BurstGPT v2.0, trace 3; n = 5,344,021 requests, no sampling. Latency is end-to-end, not TTFT.
Counts use geometric input bins and stated latency bins; quantile lines require at least 100 requests per input bin.
Zero latencies are retained (3,175,935); missing latency in traces 1 and 2 is excluded, never changed to zero.