Size and latency relationship within each LLM model

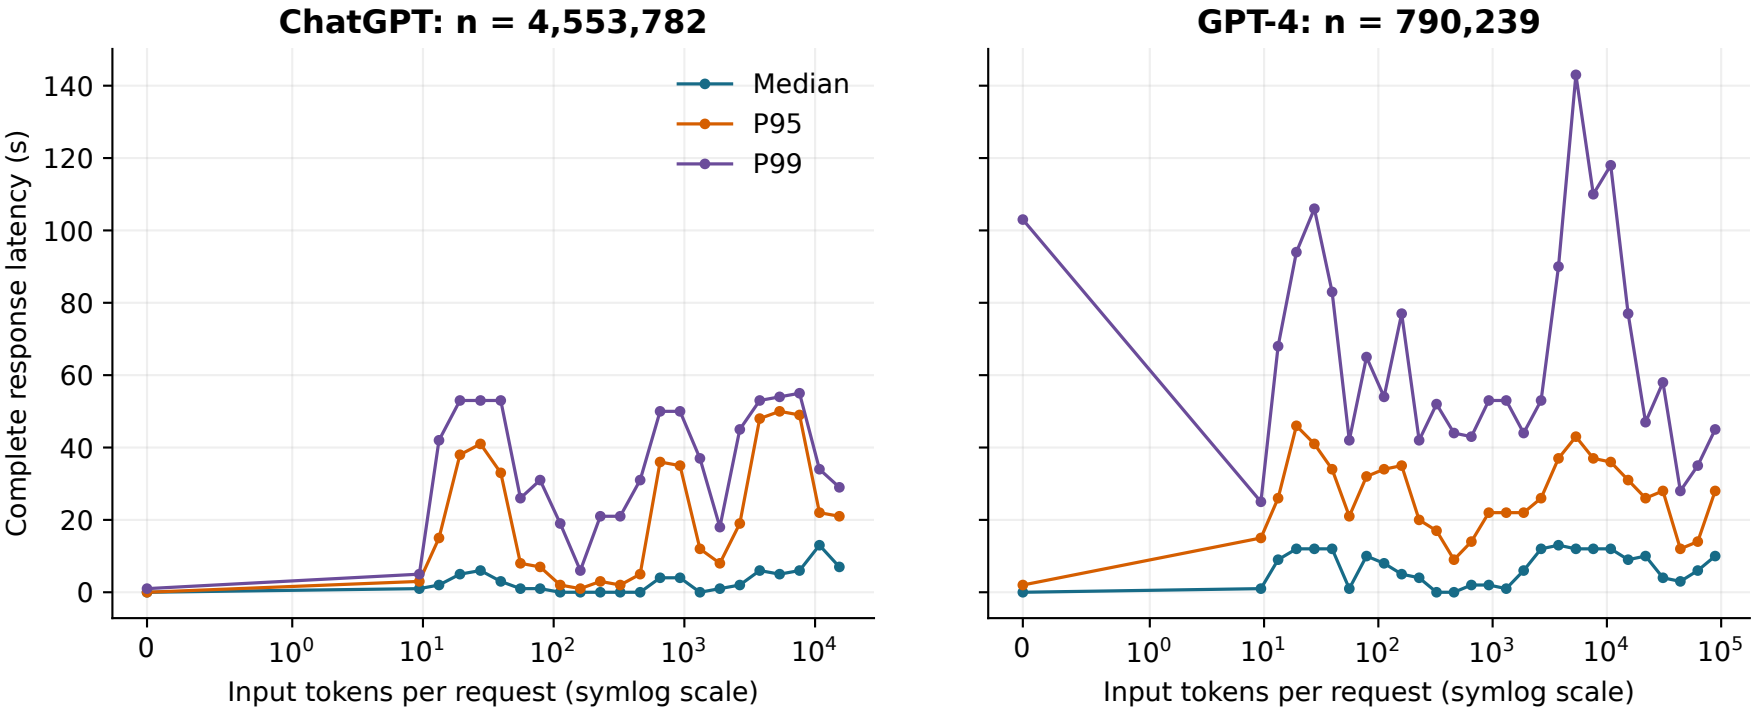

Source: HPMLL/BurstGPT v2.0, trace 3; n = 5,344,021 requests, no sampling. Exact within-bin quantiles.
Each line includes bins with at least 100 requests for that model. Zero latencies remain; hardware and service conditions are unknown.
Stratification separates the model categories but does not remove other confounding factors.