Input and output lengths have distinct, skewed distributions

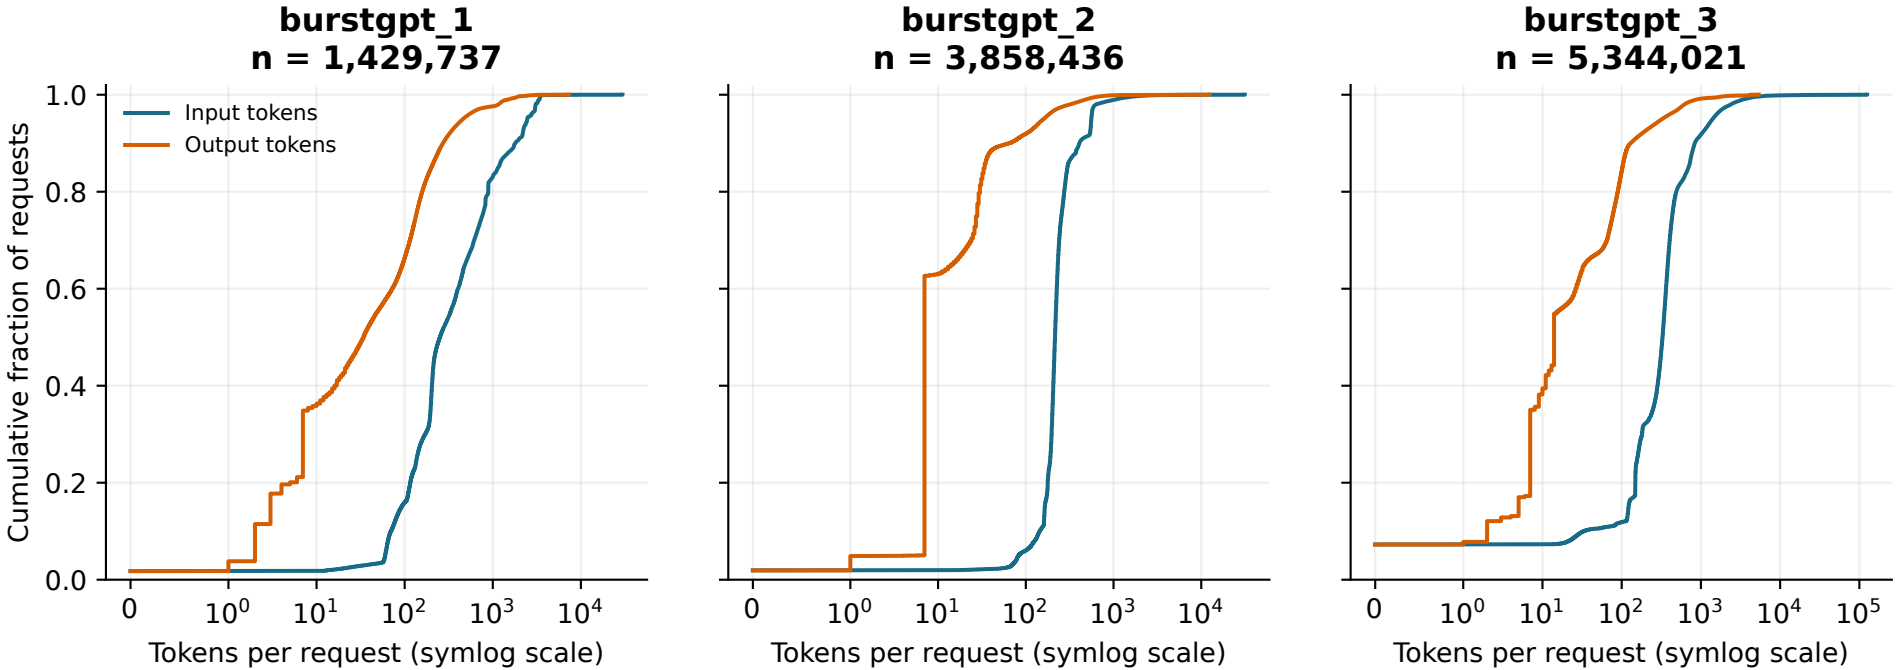

Source: HPMLL/BurstGPT v2.0; n = 10,632,194 full-trace requests. Exact empirical cumulative distributions.
Token counts are unitless; zero-token observations are retained. The linear segment around zero avoids dropping zero values.