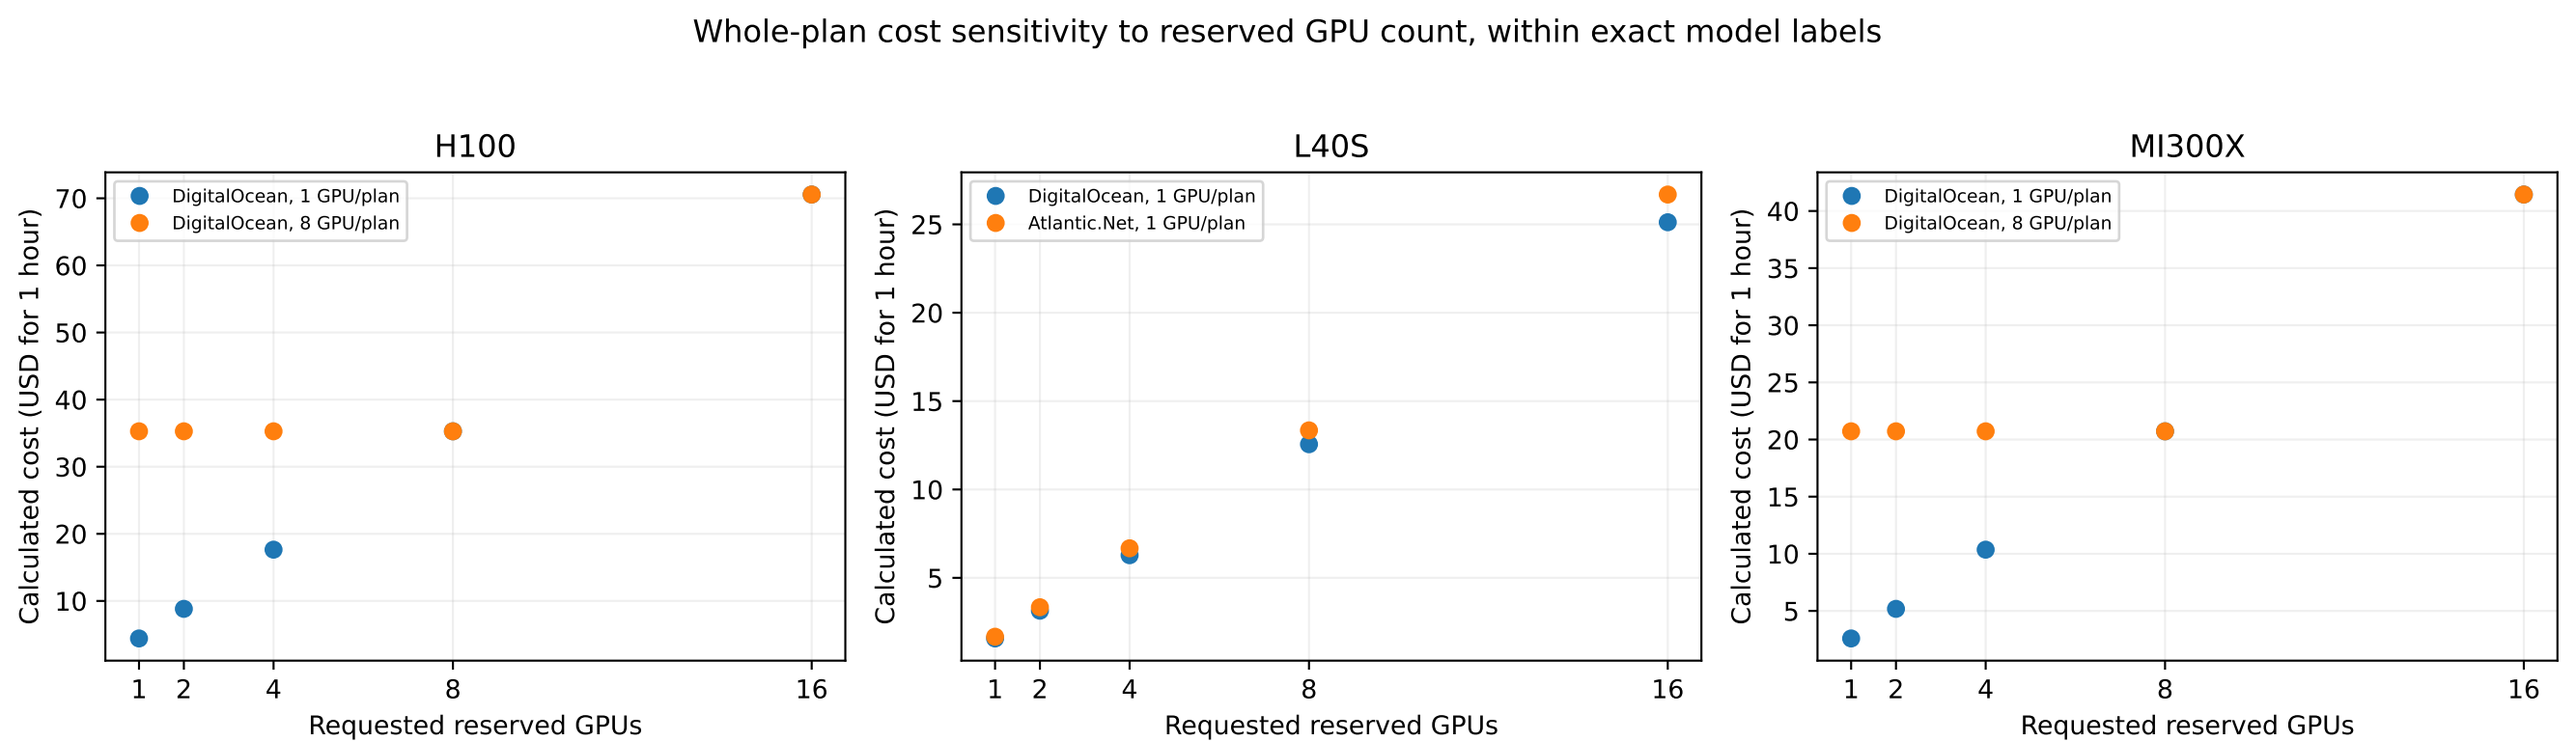

Observed inputs: GPUCloudCompare 2026-11-07; n=30 calculated offer/scenario points.
Only exact model labels shown; no cross-model speed equivalence. Points show whole-plan charges at the five stated GPU counts.
GPU counts are explicit hypothetical capacity demands; no Alibaba/BurstGPT job is assigned to these offers.